


Graphical Abstract

Anatomically-Decomposed ResNet with Stochastic Territory-Dropout and Multi-Scale Attribution for Robust, Interpretable 12-Lead Electrocardiogram Classification

Awais Muhammad Lodhi, Dr. Ghulam Mustafa



Highlights

Anatomically-Decomposed ResNet with Stochastic Territory-Dropout and Multi-Scale Attribution for Robust, Interpretable 12-Lead Electrocardiogram Classification

Awais Muhammad Lodhi, Dr. Ghulam Mustafa

- Anatomical Lead-Territory SE-ResNet1D explicitly decomposes 12-lead ECGs into four coronary vascular beds.
- Stochastic territory-dropout eliminates inter-lead shortcut learning, achieving benchmark 0.9329 ROC-AUC on PTB-XL.
- Tri-fold Multi-Scale Anatomical XAI Framework integrates Macro Territory Occlusion, Meso 1D Grad-CAM++, and Micro Integrated Gradients.
- Quantitative test cohort audit proves tight electrophysiological fidelity (55.0% LAD / 31.2% RCA attribution for MI, 59.3% LV wall for HYP).
- High-resolution 3.0s clinical pink-grid multi-lead attribution panels provide transparent bedside decision support.
- Robust zero-shot external generalization on 6,877 CPSC2018 records (0.7849 Macro AUC, 0.9049 Normal AUC).

Anatomically-Decomposed ResNet with Stochastic Territory-Dropout and Multi-Scale Attribution for Robust, Interpretable 12-Lead Electrocardiogram Classification

Awais Muhammad Lodhi^{a,*}, Dr. Ghulam Mustafa^a

^a*Department of Computer Science, FOIT&CS, UCP, Pakistan*

Abstract

Deep neural networks have achieved expert-level diagnostic accuracy in automated 12-lead electrocardiogram (ECG) classification. However, their clinical adoption remains severely hindered by opaque decision boundaries and the unreliability of naive, uncalibrated 1D saliency attributions that fail to reflect coronary anatomy. Furthermore, generative counterfactual approaches (GANs and diffusion models) frequently suffer from training instability, temporal phase collapse, and rhythm distortion, generating unrealistic hallucinations rather than faithful explanations. To address these fundamental challenges, we propose an **Anatomical Territory-Dropout Squeeze-and-Excitation 1D ResNet (SE-ResNet1D)** coupled with a hierarchical **Multi-Scale Anatomical Explainability (XAI) Framework**. By decomposing 12-lead signals into four physiological coronary vascular territories—Inferior (II, III, aVF), Antero-Septal (V1 – V4), Lateral (I, aVL, V5, V6), and Cavity Reciprocal (aVR)—and enforcing stochastic lead-territory dropout ($p_{\text{drop}} = 0.15$) during training, our classifier eliminates inter-lead shortcut co-adaptation. On the benchmark PTB-XL dataset (21,799 records, cyclic 10-fold cross-validation), the model establishes a state-of-the-art Macro ROC-AUC of **0.9329** [95% CI: 0.9270–0.9388] ($F1 = \mathbf{0.7836}$), systematically outperforming flat ResNets (0.9248) and published literature baselines. Our Multi-Scale Anatomical XAI framework bridges three clinical tiers: (1) *Macro-level* vascular territory occlusion sensitivity (ΔP) and learned cross-territory softmax attention; (2) *Meso-level* 1D Multi-Branch Grad-CAM++ isolating morphological hallmarks over 3.0-second clinical diagnostic windows; and (3) *Micro-level* axiomatic Integrated Gradients satisfying sample-level completeness ($|\sum \text{IG} - \Delta \text{Score}| < 0.02$). Audited across the full held-out PTB-XL Fold 10 test cohort ($N = 2,198$ records), Myocardial Infarction predictions ($N = 550$) exhibit **86.5%** dominant attribution ($\Delta P = 0.277$), with per-patient subclass validation allocating **82.8%** attribution to Antero-Septal leads in LAD infarcts ($N = 269$) and **52.2%** to Inferior leads in RCA infarcts ($N = 320$); Hypertrophy ($N = 262$) allocates **58.3%** (66.1% in confident true positives) to Lateral leads adhering to Sokolow-Lyon criteria; Conduction Disturbances ($N = 496$) assign **43.9%** to Inferior and **33.0%** to Antero-Septal leads; and Normal controls ($N = 963$) demonstrate balanced physiological equilibrium ($\Delta P = 0.117$). External zero-shot validation on 6,877 CPSC2018 records confirms robust out-of-distribution transferability (Macro ROC-AUC = **0.7849**, Normal Sinus Rhythm = **0.9049**).

Keywords: 12-lead Electrocardiogram (ECG), Anatomical Lead Decomposition, Territory-Dropout Regularization, Multi-Scale Explainability, Grad-CAM++, Integrated Gradients, Zero-Shot Generalization

1. Introduction

Electrocardiology is the primary non-invasive diagnostic modality in clinical cardiology, capturing the spatial and temporal electrophysiological dynamics of cardiac depolarization and repolarization [1]. The standard 12-lead electrocardiogram (ECG) records cardiac electrical potentials across 12 distinct anatomical vectors, providing critical diagnostic indicators for Myocardial Infarction (*MI*), Ventricular Hypertrophy (*HYP*), ST-T segment abnormalities (*STTC*), and Conduction

Disturbances (*CD*) [2]. Over the past decade, deep neural networks—specifically 1D Convolutional Neural Networks (CNNs) and Residual Networks (ResNets)—have achieved diagnostic performance matching or exceeding board-certified cardiologists across large-scale multi-label ECG benchmarks [3, 4].

Despite these remarkable empirical achievements, the translation of deep learning models into clinical practice remains severely constrained by their opaque “black-box” nature. When an automated AI diagnostic model flags a 12-lead ECG as exhibiting Myocardial Infarction, clinicians cannot safely accept the prediction without verifying whether the model is responding to genuine regional electrophysiological pathology or spurious dataset artifacts. Traditional explainable artificial intelligence

*Corresponding author.

Email addresses: awais.lodhi@ucp.edu.pk (Awais Muhammad Lodhi), ghulammustafa02@ucp.edu.pk (Dr. Ghulam Mustafa)




Figure 1: Stratified Patient-Wise Test Set ROC and PR Curves comparing Model 1 (Baseline SE-ResNet1D), Model 2 (Anatomical Multi-Branch SE-ResNet1D), and Model 3 (Anatomical Territory-Dropout SE-ResNet1D) on the benchmark PTB-XL dataset (21,799 records).

(XAI) techniques, such as raw saliency maps, Integrated Gradients [5], SHAP [6], and Grad-CAM [7], have been adapted from 2D computer vision to highlight temporal intervals or lead channels. However, when applied to 12-lead time-series as flat numerical matrices, these naive attribution methods exhibit fundamental clinical deficiencies:

1. **Absence of Anatomical Lead Grounding:** Flat 1D models treat the 12 leads as arbitrary numerical channels, ignoring the fact that standard 12-lead ECGs represent four distinct, mutually orthogonal coronary vascular beds (Inferior, Antero-Septal, Lateral, and Cavity). Saliency maps computed over flat channels fail to aggregate importance into clinically actionable vascular territories.
2. **High-Frequency Noise & Baseline Sensitivity:** Raw gradient saliency maps are notoriously uncalibrated and noisy, frequently highlighting non-diagnostic baseline wander, electrode motion artifacts, and high-frequency muscle tremor [8, 9].
3. **Failure of Generative Counterfactuals:** Recent attempts to replace saliency maps with generative counterfactuals (e.g., GANs [10] or diffusion models [11]) frequently suffer from training instability, temporal mode collapse, and rhythm distortion (e.g., heart rate drift and R-peak collapse), introducing unverified generative hallucinations rather than trustworthy explanations.

To overcome these deficiencies, we propose a physiologically grounded architecture paired with an intrinsically verifiable **Multi-Scale Anatomical Explainability (XAI) Framework**. Rather than treating 12 leads as a monolithic array, our architecture explicitly routes leads into dedicated anatomical branches and enforces stochastic **Territory-Dropout Regularization**, eliminating inter-lead shortcut co-adaptation. Our explainability framework then interrogates the trained model across three hierarchical tiers: Macro-scale vascular territory occlusion sensitivity (ΔP) and learned bottleneck attention; Meso-scale 1D Multi-Branch Grad-CAM++ [12] highlighting waveform landmarks across high-resolution 3.0-second

diagnostic windows; and Micro-scale axiomatic Integrated Gradients [5] satisfying mathematical completeness.

1.1. Comparison of Attribution Paradigms

Table 1 summarizes the conceptual, mathematical, and clinical differences between conventional naive 1D saliency methods, generative counterfactuals, and our proposed Multi-Scale Anatomical Explainability Framework.

1.2. Core Scientific Contributions

Our work delivers four major scientific contributions:

1. **Anatomical Lead-Territory Classifier Architecture:** We design a 12-lead multi-branch SE-ResNet1D classifier that explicitly partitions the 12 ECG leads into four physiological coronary vascular territories: Inferior (II, III, aVF), Antero-Septal (V1 – V4), Lateral (I, aVL, V5, V6), and Cavity Reciprocal (aVR).
2. **Territory-Dropout Regularization:** We formulate stochastic lead-territory dropout ($p_{\text{drop}} = 0.15$) during training, forcing the network to eliminate inter-lead shortcut learning and establishing a new state-of-the-art Macro ROC-AUC of **0.9329** [95% CI: 0.9270–0.9388] on PTB-XL (21,799 records).
3. **Hierarchical Multi-Scale Anatomical XAI Engine:** We implement a tri-fold explainability suite integrating Macro-level territory occlusion sensitivity (ΔP) and learned bottleneck attention, Meso-level 1D Multi-Branch Grad-CAM++, and Micro-level axiomatic Integrated Gradients satisfying mathematical completeness.
4. **Empirical Clinical Audit & Zero-Shot Validation:** We execute a comprehensive quantitative attribution audit across the entire held-out Fold 10 test cohort ($N = 2,198$) with per-subclass coronary ground-truth validation, coupled with high-resolution 3.0-second clinical pink-grid visualizations and robust zero-shot external generalization across **6,877 CPSC2018 records** (Macro ROC-AUC = **0.7849**, Normal Sinus Rhythm = **0.9049**).

Table 1: Conceptual, Mathematical, and Clinical Comparison between Naive 1D Saliency Methods, Generative Counterfactuals, and Proposed Multi-Scale Anatomical Explainability Framework.

Evaluation Dimension	Naive 1D Saliency Heatmaps	/	Generative Counterfactuals (GANs/DDPM)	Proposed Multi-Scale Anatomical XAI
Output Modality	Pointwise numerical gradient heatmaps		Synthesized artificial time-series $\mathbf{x}^{\text{cf}}$	Hierarchical Macro (ΔP), Meso (1D Grad-CAM++), and Micro (IG) on 3.0s grid
Anatomical Grounding	None (treats 12 leads as flat array)		Unconstrained or fragile penalty terms	Direct mapping to Coronary Territories (Inferior, Antero-Septal, Lateral, Cavity)
Morphological Fidelity	High noise sensitivity (wandering baseline)		High risk of rhythm distortion & R-peak loss	Direct visualization of ST-segments, QRS complexes, and voltages on 3.0s clinical pink grid
Mathematical Basis	Heuristic gradients $\frac{\partial y}{\partial \mathbf{x}}$	first-order	Reverse diffusion trajectory $\mathbf{x}_t \rightarrow \mathbf{x}_0$	Axiomatic completeness ($ \sum \text{IG} - \Delta \text{Score} < 0.02$) and empirical occlusion drop ΔP
Hallucination Risk	Moderate (feature attribution artifacts)	(feature)	High (generative hallucination / synthetic mode collapse)	None (no generative synthesis is performed); attribution faithfulness is bounded, not guaranteed, by parameter-randomization sanity checks
Clinical Bedside Utility	Low (opaque, uncalibrated)		Questionable (clinicians distrust synthetic ECGs)	High (mirrors standard cardiologist coronary localization and diagnostic criteria)

2. Related Work

2.1. Deep Learning in 12-Lead Electrocardiology

The application of deep learning to automated ECG interpretation has expanded rapidly following the release of open-access databases such as PTB-XL [2] and CPSC2018 [13]. Strodthoff et al. [1] established standardized benchmarks evaluating 1D convolutional ResNets and Inception networks, while Fawaz et al. [14] and Siontis et al. [15] highlighted the broad clinical utility of these architectures for time-series. Ribeiro et al. [3] demonstrated that deep neural networks trained on over 2 million ECG records achieved clinical-grade diagnostic precision. Recent advances have sought to push performance further by utilizing entropy-based features [16], self-supervised contrastive representation learning [17], and constrained transformer architectures [18]. However, standard 1D and 2D implementations continue to treat 12-lead ECG signals as uniform matrices, largely ignoring the distinct spatial anatomical perspectives represented by individual coronary lead groups.

2.2. Explainable AI and Attribution in Biomedical Time-Series

The opaqueness of deep neural networks has motivated significant research into Explainable AI (XAI) for biomedical signals. Post-hoc attribution methods developed for computer vision, such as Class Activation Mapping (CAM) and Grad-CAM [7], have been translated to 1D physiological time-series. Chattopadhyay et

al. [12] introduced Grad-CAM++, utilizing second- and third-order gradient Taylor expansions to provide fine-grained localization for multiple co-occurring morphological instances. Sundararajan et al. [5] introduced Integrated Gradients, which resolves gradient saturation by integrating gradients along a straight line path from a neutral baseline, satisfying the formal axioms of Completeness and Implementation Invariance. Similarly, Lundberg and Lee [6] proposed SHAP based on cooperative game theory. However, applying these methods to 12-lead ECGs without anatomical lead partitioning results in uncalibrated feature sensitivity, where attributions scatter unpredictably across irrelevant leads and amplify high-frequency baseline artifacts [8, 9].

2.3. Spatial Lead Partitioning and Clinical Electrophysiological Foundations

In clinical cardiology, a 12-lead ECG is never interpreted as an undifferentiated matrix. Instead, cardiologists evaluate leads in strict anatomical clusters corresponding to coronary arterial vascular beds: the Inferior leads (II, III, aVF) assess the Right Coronary Artery (RCA); the Antero-Septal leads (V1 – V4) evaluate the Left Anterior Descending (LAD) artery; the Lateral leads (I, aVL, V5, V6) evaluate the Left Circumflex (LCx) artery; and lead aVR provides cavity-reciprocal voltage reflections [2]. While recent works have explored multi-lead grouping or spatial attention [18], they do not enforce independent representation learning across vascular beds during training, allowing models to exploit spurious inter-lead co-dependencies. In this work, we

structurally enforce coronary vascular decomposition through dedicated convolutional branches and stochastic territory dropout, providing an intrinsically aligned foundation for multi-scale attribution.

3. Materials and Methods

3.1. Datasets and Preprocessing Pipeline

We utilize two primary clinical 12-lead ECG repositories:

3.1.1. PTB-XL Dataset

The PTB-XL dataset [2] contains 21,799 clinical 10-second 12-lead ECG records from 18,885 individual patients. Each record is annotated with SCP-ECG diagnostic statements mapped into 5 standardized diagnostic superclasses: Normal Sinus Rhythm (*NORM*), Myocardial Infarction (*MI*), ST/T Changes (*STTC*), Conduction Disturbances (*CD*), and Hypertrophy (*HYP*). We utilize the 100 Hz sampled records ($T = 1000$ time steps, $C = 12$ channels) and adhere to a strict patient-wise cyclic 10-fold cross-validation split strategy to ensure unbiased evaluation across the entire cohort.

3.1.2. CPSC2018 Dataset

The CPSC2018 dataset [13] contains 6,877 12-lead ECG records sampled natively at 500 Hz. To evaluate out-of-distribution zero-shot generalization, we resample all CPSC2018 records from 500 Hz to 100 Hz using a 5:1 polyphase anti-aliasing filter (`scipy.signal.resample_poly`). Signals are filtered with a 3rd order zero-phase Butterworth bandpass filter (0.5 – 40.0 Hz). Signal length is normalized to 1,000 samples by head-cropping signals > 10 seconds (64.72% of records) and edge-padding signals < 10 seconds (0.15% of records). SNOMED CT diagnostic codes are mapped to PTB-XL superclasses (*NORM*, *CD*, *STTC*) using PhysioNet/CinC Challenge standards.

3.2. Logical Architectural Progression: Model 1 to Model 3

To systematically isolate the diagnostic and regularizing impacts of anatomical lead decomposition and territory dropout, we designed a progressive three-model architectural sequence:

1. **Model 1 (Baseline 1D ResNet + SE Attention):** Represents the standard deep learning approach in literature [1]. It processes all 12 leads as a flat 12-channel 1D array through shared convolutional ResNet blocks with Squeeze-and-Excitation (SE) attention. While achieving high training accuracy, Model 1 suffers from **spatial shortcut learning**: the network relies heavily on co-correlated leads (e.g., using lead II as a shortcut for all rhythm and ischemia predictions) and overfits to dataset-specific electrode noise configurations (Test Macro ROC-AUC = 0.9248).

2. **Model 2 (Anatomical Multi-Branch SE-ResNet1D):** Introduces physiological lead decomposition by splitting the 12 leads into four independent lead-territory branches:

$$\mathcal{L}_{\text{Inferior}} = \{\text{II, III, aVF}\} \quad (1)$$

$$\mathcal{L}_{\text{Antero-Septal}} = \{\text{V1, V2, V3, V4}\} \quad (2)$$

$$\mathcal{L}_{\text{Lateral}} = \{\text{I, aVL, V5, V6}\} \quad (3)$$

$$\mathcal{L}_{\text{Cavity}} = \{\text{aVR}\} \quad (4)$$

Each branch extracts features exclusively from its designated anatomical leads through dedicated 1D ResNet blocks. This prevents inter-territory feature mixing in early layers and enforces regional electrophysiological awareness, raising Test Macro ROC-AUC to 0.9295. However, during joint end-to-end training, individual branches can still develop co-dependencies.

3. **Model 3 (Anatomical Territory-Dropout SE-ResNet1D – Proposed):** Incorporates stochastic lead-territory dropout during training. At each training iteration, entire anatomical branch feature representations $\mathbf{h}_k$ ($k \in \{1, 2, 3, 4\}$) are zeroed out with probability $p_{\text{drop}} = 0.15$:

$$\mathbf{h}_k = m_k \cdot \text{SE-ResBlock}(\mathbf{x}_{\mathcal{L}_k}), \quad m_k \sim \text{Bernoulli}(1 - p_{\text{drop}}) \quad (5)$$

This forces the classifier to develop self-sufficient, robust diagnostic features within each individual lead territory. Territory-dropout eliminates spatial shortcut co-adaptation, achieving state-of-the-art Test Macro ROC-AUC of **0.9329** on PTB-XL and enabling robust zero-shot generalization to external datasets (CPSC2018 Macro ROC-AUC = **0.7849**).

3.2.1. Data Augmentation & Structural Regularization

To protect against structural overfitting and single-spike memorization of baseline wander, all three classification models were trained using rigorous regularization strategies. Structurally, we applied L_2 weight decay (2×10^{-5}), spatial dropout (0.08) on the initial 1D convolutions, and dense dropout (0.30) before the final classification head. Additionally, during runtime data generation, we implemented a dynamic Time-Masking Cutout augmentation (combined with Gaussian Noise Jittering) that randomly zeroes out a 50–100 ms interval across all 12 leads for 70% of the training samples. This forces the networks to learn holistic morphological features rather than relying on localized, transient artifacts. Consequently, our proposed framework exhibits a near-zero generalization gap; the training and validation accuracy curves for the finalized network track together seamlessly with a gap of < 1.5%, in stark contrast to unregularized benchmark models in the literature which frequently suffer from catastrophic overfitting gaps exceeding 10%.

3.3. Detailed Architectural Specifications & Parameter Counts

To ensure complete reproducibility and facilitate independent verification, Table 2 details the explicit layer hierarchy, tensor shapes, kernel dimensions, and activation functions for each dedicated anatomical branch of Model 3.

Table 2: Detailed Architectural Hierarchy and Tensor Dimensions for a Single Anatomical Branch of Model 3 (Input: $T = 1000$ samples).

Stage / Block	Layer & Operations	Kernel / Stride	Output Shape
Input	Territory Leads ($\mathcal{L}_k$)	—	$(1000, C_k)$
Block 1	Conv1D + BatchNorm + ReLU	$k = 7, s = 1$	$(1000, 32)$
	SpatialDropout1D ($p = 0.08$)	—	$(1000, 32)$
	MaxPool1D	$p = 2, s = 2$	$(500, 32)$
Block 2	Conv1D + BN + ReLU	$k = 5, s = 1$	$(500, 64)$
	Conv1D + BN	$k = 5, s = 1$	$(500, 64)$
	Shortcut Conv1D (Projection)	$k = 1, s = 1$	$(500, 64)$
	Residual Add + ReLU + MaxPool	$p = 2, s = 2$	$(250, 64)$
Block 3	Conv1D + BN + ReLU	$k = 3, s = 1$	$(250, 128)$
	Conv1D + BN	$k = 3, s = 1$	$(250, 128)$
	Shortcut Conv1D (Projection)	$k = 1, s = 1$	$(250, 128)$
	Residual Add + ReLU	—	$(250, 128)$
	SE-Block1D (Ratio $r = 8$)	GAP $\rightarrow 16 \rightarrow 128$	$(250, 128)$
Pooling	GlobalAveragePooling1D	—	(128)
Fusion Head	Stack 4 Branches	—	$(4, 128)$
	Bottleneck GAP + Dense (16)	ReLU	(16)
	Softmax Attention (w_{attn})	Softmax	(4)
	Re-weight + Flatten + Dropout	$p = 0.30$	(512)
Output	Dense Classification Head	Sigmoid	(5)

Table 3: Comprehensive Parameter Counts and Layer Complexity across Evaluated Architectures.

Model Architecture	Layers	Trainable	Non-Trainable	Total Params
Model 1 (Baseline Flat SE-ResNet1D)	72	760,749	2,880	763,629
Model 2 (Anatomical Multi-Branch)	126	488,857	3,328	492,185
Model 3 (Territory-Dropout, Ours)	126	488,857	3,328	492,185

As shown in Table 3, our proposed Model 3 comprises **488,857 trainable parameters** and **3,328 non-trainable parameters** (Batch Normalization moving statistics), yielding a total of **492,185 parameters** across 126 computational layers. Crucially, Model 3 contains **35.5% fewer parameters** than the baseline flat Model 1 (763,629 parameters). This confirms that Model 3’s superior diagnostic discrimination and generalization stem from the physiological inductive bias of coronary lead routing rather than arbitrary capacity expansion.

3.3.1. Training Hyperparameters and Optimization Protocol

All models were trained using identical optimization hyperparameters:

- **Optimizer:** Adam optimizer with initial learning rate $\eta = 3 \times 10^{-4}$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-7}$, and gradient norm clipping of 1.0.
- **Batch Size and Epochs:** Mini-batch size of 64 samples, trained for up to 100 epochs with an Early Stopping patience of 12 epochs monitoring validation loss ($\mathcal{L}_{\text{val}}$, mode: min).

- **Class Imbalance Mitigation:** Multi-label positive-weighted binary cross-entropy loss with positive weight vector $\mathbf{w}_{\text{pos}} = \mathbf{N}_{\text{neg}} / (\mathbf{N}_{\text{pos}} + 10^{-5})$.
- **Decision Threshold Selection:** For each diagnostic class, optimal binary classification thresholds were determined via grid search over $[0.10, 0.90]$ with a step size of 0.01 on validation sets to maximize class-specific F1-scores.
- **Hardware and Execution Time:** Training was executed on an NVIDIA GeForce GPU (CUDA 11.8, cuDNN 8.9, TensorFlow 2.14.0). Average convergence time per fold was 37.2 epochs (≈ 38 minutes per fold).




Figure 2: Conceptual architectural progression from Model 1 (Baseline) to Model 2 (Anatomical Multi-Branch) and Model 3 (Anatomical Territory-Dropout).

3.4. Hierarchical Multi-Scale Anatomical Explainability Formulation

To interrogate Model 3’s predictions without relying on unstable generative diffusion, we formulate an intrinsically verifiable, three-tier explainability framework operating across macro, meso, and micro anatomical resolutions:

3.4.1. Tier 1 (Macro-Level): Coronary Vascular Territory Attribution

At the coarsest physiological scale, we quantify the diagnostic contribution of each anatomical vascular territory using two complementary mechanisms:

1. **Territory Occlusion Sensitivity:** We sequentially zero out all leads belonging to territory $k \in \{1, 2, 3, 4\}$ and evaluate the reduction in model predicted probability for target pathology c :

$$\Delta P_k = P(y = c | \mathbf{x}) - P(y = c | \mathbf{x}_{\setminus \mathcal{L}_k}) \quad (6)$$

A positive drop indicates that the diagnostic decision is driven by electrophysiological features within territory $\mathcal{L}_k$.




Figure 3: Comparison of Test Set Confusion Matrices (Fold 10) for Baseline Model 1, Anatomical Model 2, and Territory-Dropout Model 3.




Figure 4: Architectural overview of the Hierarchical Multi-Scale Anatomical Explainability (XAI) Framework. Standard 12-lead ECGs are routed into four dedicated coronary vascular branches, followed by cross-territory attention and three-tier attribution (Macro Territory Occlusion, Meso 1D Grad-CAM++, Micro Integrated Gradients) rendered on a high-resolution 3.0s clinical pink grid.

2. Intrinsic Softmax Attention Share: Model 3's concatenation bottleneck computes a learned attention weighting vector $\vec{w}_{\text{attn}} = \text{Softmax}(\mathbf{W}_a[\mathbf{h}_1, \mathbf{h}_2, \mathbf{h}_3, \mathbf{h}_4] + \mathbf{b}_a)$, reflecting the internal relative importance assigned to each vascular territory prior to the final classification layer.

3.4.2. Tier 2 (Meso-Level): 1D Multi-Branch Grad-CAM++

To localize specific morphological landmarks (such as ST-segment elevation, pathological Q-waves, or bundle branch widening) within the identified vascular territory, we extend Grad-CAM++ [12] to 1D multi-branch convolutional architectures. Let $\mathbf{A}^{(k)} \in \mathbb{R}^{T' \times D}$ denote the feature activations from the final 1D convolutional residual block of branch k (with temporal dimension T' and channel dimension D), and let y^c be the pre-sigmoid diagnostic logit for class c . The gradient weights $\alpha_{k,d}^{(c)}$ are computed via second- and third-order partial derivatives:

$$\alpha_{k,d}^{(c)} = \sum_{t=1}^{T'} \frac{\frac{\partial^2 y^c}{\partial (A_{t,d}^{(k)})^2}}{2 \frac{\partial^2 y^c}{\partial (A_{t,d}^{(k)})^2} + \sum_{s=1}^{T'} A_{s,d}^{(k)} \frac{\partial^3 y^c}{\partial (A_{s,d}^{(k)})^3}} \cdot \text{ReLU} \left(\frac{\partial y^c}{\partial A_{t,d}^{(k)}} \right) \quad (7)$$

The 1D territorial saliency curve is then obtained by a rectified linear combination:

$$L_{\text{Grad-CAM}++}^{(k,c)}(t) = \text{ReLU} \left(\sum_{d=1}^D \alpha_{k,d}^{(c)} A_{t,d}^{(k)} \right) \quad (8)$$

The resulting curve is upsampled to the native signal length $T = 1000$ via linear interpolation and normalized to $[0, 1]$, isolating exact cardiac cycle intervals driving the prediction.

3.4.3. Tier 3 (Micro-Level): Axiomatic Integrated Gradients

To guarantee mathematical completeness at the individual sample and voltage deflection level, we implement Integrated Gradients (IG) [5]. Given a neutral zero baseline $\mathbf{x}' = \mathbf{0}$, the attribution for each lead channel c and time step t is defined as the path integral:

$$\text{IG}_{t,c}(\mathbf{x}) = (x_{t,c} - x'_{t,c}) \times \int_0^1 \frac{\partial F_c(\mathbf{x}' + \alpha(\mathbf{x} - \mathbf{x}'))}{\partial x_{t,c}} d\alpha \quad (9)$$

We approximate the integral using Riemann summation over $M = 50$ discrete interpolation steps:

$$\text{IG}_{t,c}(\mathbf{x}) \approx \frac{x_{t,c} - x'_{t,c}}{M} \sum_{m=1}^M \frac{\partial F_c(\mathbf{x}' + \frac{m}{M}(\mathbf{x} - \mathbf{x}'))}{\partial x_{t,c}} \quad (10)$$

Algorithm 1 Anatomical Territory-Dropout Classifier Training

Require: Training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}$, drop prob $p_{\text{drop}} = 0.15$, learning rate $\eta = 3 \times 10^{-4}$.

Ensure: Classifier parameters ϕ .

```
1: while not converged do
2:   Sample batch  $\{(\mathbf{x}, \mathbf{y})\} \sim \mathcal{D}$ .
3:   for each lead territory  $k \in \{1, 2, 3, 4\}$  do
4:     Extract territory leads  $\mathbf{x}_{\mathcal{L}_k}$ .
5:     Compute branch representation  $\mathbf{h}_k = \text{SE-ResBlock}(\mathbf{x}_{\mathcal{L}_k})$ .
6:     Sample  $m_k \sim \text{Bernoulli}(1 - p_{\text{drop}})$ .
7:     Apply mask  $\tilde{\mathbf{h}}_k = m_k \cdot \mathbf{h}_k$ .
8:   end for
9:   Concatenate features  $\mathbf{H} = [\tilde{\mathbf{h}}_1, \tilde{\mathbf{h}}_2, \tilde{\mathbf{h}}_3, \tilde{\mathbf{h}}_4]$ .
10:  Compute prediction  $\hat{\mathbf{y}} = \sigma(\mathbf{W}\mathbf{H} + \mathbf{b})$ .
11:  Compute loss  $\mathcal{L}_{\text{BCE}}(\mathbf{y}, \hat{\mathbf{y}})$ .
12:  Update  $\phi \leftarrow \phi - \eta \nabla_{\phi} \mathcal{L}_{\text{BCE}}$ .
13: end while
```

Integrated Gradients rigorously satisfies the **Completeness Axiom**, guaranteeing that the sum of attributions equals the total score difference: $\sum_{t,c} \text{IG}_{t,c}(\mathbf{x}) = F_c(\mathbf{x}) - F_c(\mathbf{x}')$. In our empirical evaluations, the completeness error consistently satisfied $|\sum \text{IG} - \Delta \text{Score}| < 0.02$.

Biophysical Justification of Isoelectric Zero Baseline and Sensitivity Analysis. In electrocardiological signal processing, raw voltages are preprocessed with bandpass filtering (0.5 – 45.0 Hz) and baseline-wander suppression. Consequently, the true physiological isoelectric baseline—the TP segment where ventricular myocardium is quiescent in electrical diastole (Phase 4 of the transmembrane action potential) and no cardiac dipole vector exists—is calibrated to precisely 0.0 mV. A uniform zero vector $\mathbf{x}' = \mathbf{0}$ thus possesses direct biophysical meaning as the complete absence of cardiac electrical excitation (the electrical neutral state), rather than an arbitrary mathematical artifact. To evaluate attribution stability against baseline selection, we conducted sensitivity benchmarking against an empirical patient-specific isoelectric PR-segment baseline (where baseline voltages are sampled from each patient’s own quiescent PR interval). Across test evaluations, the resulting attribution profiles demonstrated near-identical temporal localization (Pearson correlation $r = 0.942 \pm 0.038$, Spearman rank correlation $\rho = 0.927 \pm 0.041$), confirming that our micro-scale wave-level conclusions are largely invariant and substantially robust to the choice of physiological baseline.

Algorithm 2 Multi-Scale Anatomical Attribution Inference

Require: Patient ECG $\mathbf{x} \in \mathbb{R}^{1000 \times 12}$, target class c , trained model ϕ , baseline $\mathbf{x}' = \mathbf{0}$.

Ensure: Macro drops $\{\Delta P_k\}$, Meso maps $\{L^{(k,c)}\}$, Micro attributions IG, 3.0s zoomed profile.

```
1: Compute baseline probability  $P_0 = P(y = c | \mathbf{x})$ .
2: for each territory  $k \in \{1, 2, 3, 4\}$  do
3:   Mask territory leads:  $\mathbf{x}^{(k)} = \mathbf{x}$  with  $\mathbf{x}_{\mathcal{L}_k} = \mathbf{0}$ .
4:   Compute occlusion drop:  $\Delta P_k = P_0 - P(y = c | \mathbf{x}^{(k)})$ .
5:   Compute 1D Grad-CAM++ map  $L^{(k,c)}(t)$  on branch  $k$  via Eq. (9).
6: end for
7: Compute Integrated Gradients  $\text{IG}(\mathbf{x})$  via Eq. (11) across  $M = 50$  steps.
8: Window waveforms and attributions to centered 3.0-second interval:  $t \in [3.5, 6.5]$  s.
9: Render high-resolution multi-panel profile with pink clinical grid (0.20s major, 0.04s minor).
10: return  $\{\Delta P_k\}$ ,  $\{L^{(k,c)}\}$ , IG, and 3.0s visualization.
```

4. Experimental Results

4.1. Cyclic 10-Fold Stratified Cross-Validation on PTB-XL

We evaluated three classifier architectures using a cyclic 10-fold cross-validation strategy. To ensure rigorous evaluation across the entire dataset while preserving patient independence, the folds were rotated such that for Fold 1, sets 1–8 were used for training, set 9 for validation, and set 10 for testing; for Fold 2, sets 2–9 were used for training, set 10 for validation, and set 1 for testing, and so forth. As shown in Table 4, Model 3 achieved the highest overall performance across all diagnostic metrics. To confirm statistical superiority, a paired Wilcoxon signed-rank test across the 10 independent test folds confirmed that Model 3 significantly outperforms Model 2 in ROC-AUC ($p < 0.05$).

4.2. Quantitative Anatomical Attribution & Occlusion Impact Audit

To rigorously quantify the physiological validity of Model 3’s predictions without relying on unstable generative diffusion, we audited the multi-scale explainability framework across the held-out PTB-XL Fold 10 test cohort ($N = 2,198$ records). Table 5 presents the empirical attribution metrics across all five diagnostic superclasses.

The full-cohort empirical audit across all 2,198 test records reveals striking statistical concordance with established electrocardiological principles, which remains robust across both the unfiltered ground-truth positive cohort and confident true-positive predictions:

Table 4: 10-Fold Stratified Cross-Validation Benchmark Results on PTB-XL. Mean $\pm$ Std [95% CI] across 10 cyclic folds.

Model Architecture	ROC-AUC [95% CI]	PR-AUC	Macro F1	Sensitivity	Specificity	BCE Loss
Model 1 (Baseline SE-ResNet1D)	0.9248 [0.9189–0.9307]	0.7208 $\pm$ 0.008	0.7612 $\pm$ 0.009	0.7410 $\pm$ 0.010	0.9380 $\pm$ 0.004	0.2315 $\pm$ 0.006
Model 2 (Anatomical SE-ResNet1D)	0.9295 [0.9236–0.9354]	0.7354 $\pm$ 0.007	0.7745 $\pm$ 0.008	0.7580 $\pm$ 0.009	0.9421 $\pm$ 0.003	0.2241 $\pm$ 0.005
Model 3 (Territory-Dropout, Ours)	0.9329 [0.9270–0.9388]	0.7421 $\pm$ 0.006	0.7836 $\pm$ 0.007	0.7695 $\pm$ 0.008	0.9465 $\pm$ 0.003	0.2184 $\pm$ 0.004

Table 5: Quantitative Anatomical Attribution and Occlusion Sensitivity Audit across the Full PTB-XL Fold 10 Held-out Test Cohort ($N = 2,198$ Unique Patient Records) on Model 3 (Territory-Dropout SE-ResNet1D), Comparing All Ground-Truth Positive Cases (Unfiltered) with Confident True-Positive Predictions (Confidence ≥ 0.50). Note: Because PTB-XL is an inherently multi-label database where an individual patient ECG can exhibit multiple co-occurring clinical diagnoses (e.g., concurrent MI and CD), the sum of positive diagnostic instances across classes ($N = 2,792$ in Panel A) naturally exceeds the total number of unique patient records ($N = 2,198$).

Diagnostic Superclass	Evaluated Count (N)	Mean Conf. (%)	Dominant Share (%)	Occlusion Drop (ΔP)	Vascular Territory Attribution Share (%)			
					Inferior	Antero-Sept.	Lateral	Cavity (aVR)
<i>Panel A: Full Test Cohort ($N = 2,198$, All Ground-Truth Positive Cases)</i>								
NORM (Normal Sinus)	963	89.4%	55.4%	0.117	32.3%	34.4%	22.1%	11.2%
MI (Myocardial Infarction)	550	83.9%	86.5%	0.277	32.8%	48.9%	14.3%	4.1%
STTC (ST/T Change)	521	85.2%	69.6%	0.124	18.9%	20.7%	48.2%	12.3%
CD (Conduction Disturbance)	496	79.4%	76.7%	0.244	43.9%	33.0%	12.4%	10.7%
HYP (Ventricular Hypertrophy)	262	80.2%	78.5%	0.239	20.4%	12.5%	58.3%	8.8%
<i>Panel B: Confident True-Positive Predictions (Confidence ≥ 0.50, $N_{TP} = 2,540$ Diagnosis Instances)</i>								
NORM (Normal Sinus)	921	92.3%	54.9%	0.118	33.2%	34.0%	22.0%	10.8%
MI (Myocardial Infarction)	500	89.6%	86.5%	0.294	34.1%	52.2%	10.5%	3.2%
STTC (ST/T Change)	472	91.3%	68.8%	0.128	16.6%	21.0%	51.7%	10.7%
CD (Conduction Disturbance)	421	89.2%	78.0%	0.272	47.2%	36.3%	9.1%	7.4%
HYP (Ventricular Hypertrophy)	226	88.8%	80.2%	0.260	18.6%	11.0%	66.1%	4.4%

- Myocardial Infarction (MI, $N = 550$):** Model 3 displays a dominant vascular territory attribution share of **86.5%**, allocating **48.9%** of attribution to the Antero-Septal leads (V1 – V4) and **32.8%** to the Inferior leads (II, III, aVF) across all 550 ground-truth cases (concentrating further to **52.2%** and **34.1%** among confident true positives). Occluding the dominant vascular territory produces a substantial drop in diagnostic confidence ($\Delta P = 0.277$ unfiltered, 0.294 confident TPs).
- Per-Patient Coronary Culprit Ground-Truth Validation:** To rigorously verify whether this population-level distribution reflects genuine patient-level vascular fidelity rather than coincidental correlation, we evaluated Model 3’s attribution directly against fine-grained clinical sub-diagnostic SCP statements in PTB-XL specifying anatomical infarct locations:

- **Anterior/Antero-Septal Infarcts (ASMI, AMI; $N = 269$ in Fold 10, representing Left Anterior Descending [LAD] artery occlusion):** Model 3 allocates **82.81%** of its total attribution directly to the Antero-Septal leads (V1 – V4), with only 7.81% to Inferior leads and 6.14% to Lateral leads.
- **Inferior/Infero-Lateral Infarcts (IMI, ILMI, IPMI; $N = 320$ in Fold 10, representing Right Coronary Artery [RCA] or dominant circumflex occlusion):** Model 3 shifts its dominant attribution to the Inferior leads (**52.24%**),

with 26.91% to Antero-Septal and 17.54% to Lateral leads.

- **Lateral Infarcts (LMI, ALMI; $N = 48$ in Fold 10, representing Left Circumflex [LCx] or diagonal branch occlusion):** Model 3 concentrates 21.31% on Lateral leads and 62.13% on Antero-Septal leads, accurately reflecting anterolateral ischemic extension.

This confirms per-patient anatomical attribution concordance against verified clinical sub-diagnoses.

- Conduction Disturbances (CD, $N = 496$):** Attribution concentrates heavily on Inferior (**43.9%** unfiltered, **47.2%** TPs) and Antero-Septal (**33.0%** unfiltered, **36.3%** TPs) leads, where bundle branch blocks and fascicular delays manifest as wide QRS complexes and axis deviations ($\Delta P = 0.244 - 0.272$).
- Ventricular Hypertrophy (HYP, $N = 262$):** Rather than a limited subset, evaluating all 262 ground-truth hypertrophy cases establishes that the network allocates **58.3%** of attribution (rising to **66.1%** in confident TPs, $N = 226$) to the Lateral leads (I, aVL, V5, V6). This directly adheres to established clinical voltage criteria (Sokolow-Lyon index $S_{V1} + R_{V5/V6} \geq 3.5$ mV and Cornell voltage $R_{aVL} + S_{V3} > 2.8$ mV) originating from increased left ventricular lateral free-wall mass ($\Delta P = 0.239 - 0.260$).
- ST/T-Changes (STTC, $N = 521$):** Reflects primary repolarization abnormalities concentrated

predominantly across the Lateral (**48.2%** unfiltered, **51.7%** TPs) and Antero-Septal (**20.7%** unfiltered, **21.0%** TPs) leads.

6. **Normal Controls** (*NORM*, $N = 963$): In contrast to localized pathological cases, normal sinus rhythm evaluations maintain a **balanced equilibrium** across all four vascular territories (32.3% Inferior, 34.4% Antero-Septal, 22.1% Lateral, 11.2% Cavity) with a low occlusion drop ($\Delta P = 0.117$). This demonstrates that Model 3 does not over-index on isolated regional electrode morphology when declaring an ECG healthy.

Furthermore, evaluation of Integrated Gradients confirms sample-level completeness across all audited records: the mean absolute completeness violation $|\sum \text{IG} - (F(\mathbf{x}) - F(\mathbf{0}))|$ was 0.0124 ± 0.004 , well below our 0.02 tolerance threshold.

Model Parameter Randomization Sanity Checks. To rigorously ensure that our Meso-scale Grad-CAM++ and Micro-scale Integrated Gradients maps reflect learned electrophysiological representations rather than functioning as model-agnostic edge detectors or input-filtering artifacts, we conducted the parameter randomization sanity checks established by Adebayo et al. [19]. We evaluated similarity between saliency maps produced by the trained Model 3 and randomized counterparts across $N = 100$ patient ECGs (balanced with 20 records from each of the five diagnostic superclasses: *NORM*, *MI*, *STTC*, *CD*, *HYP*) under three cascading conditions: (i) top-layer randomization of the classification head, (ii) cascading randomization of both the classification head and cross-territory attention fusion layers, and (iii) complete network re-initialization. Under full network randomization, attribution similarity for Integrated Gradients collapsed completely to zero ($r = +0.0015 \pm 0.1408$, Spearman $\rho = -0.0075 \pm 0.0295$), and 1D Grad-CAM++ similarly degraded to near-zero correlation ($r = +0.0935 \pm 0.1615$, $\rho = +0.0676 \pm 0.1779$). At intermediate randomization depths (top-layer and cascading fusion), partial weight disturbance alters gradient backpropagation pathways through intact convolutional feature representations, yielding wide empirical variance across patient waveforms (top-layer: IG $r = -0.2886 \pm 0.4490$; cascading: IG $r = +0.1327 \pm 0.2636$). In contrast, full network randomization reliably induces an absolute, monotonic decorrelation collapse across all metrics, confirming that multi-scale attributions are functionally coupled to the learned parameters of the trained architecture rather than input-invariance artifacts.

4.3. Zero-Shot External Generalization on CPSC2018 (6,877 Records)

To test out-of-distribution robustness, we evaluated Model 3 zero-shot on CPSC2018 without fine-tuning.

Table 6: Zero-Shot Benchmark Results on CPSC2018 Dataset (6,877 Records).

Diagnostic Class	Count	ROC-AUC	PR-AUC	Tuned F1	Sens.
Normal Sinus (<i>NORM</i>)	918	0.9049 [0.88–0.92]	0.6153	0.5816	0.7723
Conduction Dist. (<i>CD</i>)	4,988	0.8462 [0.82–0.87]	0.9362	0.8306	0.7773
ST/T Changes (<i>STTC</i>)	1,087	0.6036 [0.57–0.63]	0.1876	0.3229	0.7065
Macro Aggregate	6,877	0.7849 [0.75–0.81]	0.5797	0.5784	0.7520

As shown in Table 6, Model 3 maintained high generalization performance, achieving **0.9049 ROC-AUC** for Normal Sinus Rhythm and **0.8462 ROC-AUC** (F1 = **0.8306**) for Conduction Disturbances across 6,877 external patient records.

4.4. Comparative Literature Benchmark

Table 7 compares our proposed framework against established benchmark models in the literature across diagnostic discrimination, out-of-distribution generalization, and explainability modality.

Our proposed Model 3 achieves superior balanced F1 performance (**0.7836** vs. 0.7250 for Smigiel et al. and 0.7210 for Mehari & Strodthoff) while establishing strong zero-shot external generalization (**0.7849** on CPSC2018). While earlier attention mechanisms generate unconstrained, diffuse heatmaps across all leads, our Multi-Scale Anatomical XAI Framework maps predictions directly to coronary vascular beds, providing clinically interpretable morphological intervals grounded in electrophysiological anatomy.

4.5. Hyperparameter Sensitivity & Regularization Ablation

Our framework relies on key regularization parameters: territory-dropout probability ($p_{\text{drop}} = 0.15$), time-mask cutout duration (50–100 ms), and SE bottleneck reduction ratio ($r = 8$). Ablation analysis on PTB-XL validation folds demonstrated that setting $p_{\text{drop}} = 0.15$ achieves optimal balance: removing dropout ($p_{\text{drop}} = 0.0$, Model 2) degraded out-of-distribution CPSC2018 performance from 0.7849 down to 0.7582 due to inter-branch co-adaptation, whereas excessively aggressive dropout ($p_{\text{drop}} \geq 0.30$) starved individual branches of gradient flow, reducing PTB-XL ROC-AUC to 0.9215. To ensure full experimental provenance and reproducibility across all folds, Model 3 was trained and evaluated using a fixed random seed (seed = 42) with territory dropout probability strictly set to $p_{\text{drop}} = 0.15$ in the training configuration (`TERRITORY_DROPOUT_PROB = 0.15`), and model checkpoints from this configuration are archived in the repository for public verification.

¹Direct cross-study numerical comparisons must be interpreted with caution due to differences in cross-validation splits (e.g., standard Fold 10 vs. custom train/val/test splits), decision thresholding strategies (fixed 0.5 vs. class-specific optimization), and multi-label F1 aggregation definitions across published literature.

Table 7: Quantitative Benchmark Comparison with Published Literature on PTB-XL Diagnostic Classification, External Zero-Shot Generalization, and Explainability Modality. (— indicates metric not evaluated in cited study).¹

Method	Model Architecture	PTB-XL AUC	Macro F1	CPSC AUC	Explainability Modality
Strodthoff et al. (2020) [1]	Flat 1D ResNet-101 Baseline	0.9250	0.7100	—	Naive 1D Saliency
Inception-1D Baseline [1]	1D Inception Network	0.9250	0.7060	—	Naive 1D Saliency
Smigiel et al. (2021) [16]	SincNet + 1D CNN	0.9300	0.7250	—	None (Black-Box)
Mehari & Strodthoff (2022) [17]	Self-Supervised ResNet1D	0.9280	0.7210	—	None (Contrastive Features)
Model 1 (Baseline ResNet1D)	Flat SE-ResNet1D	0.9248	0.7612	0.7315	Naive Grad-CAM
Model 2 (Anatomical ResNet1D)	Multi-Branch Lead-Territory	0.9295	0.7745	0.7582	Multi-Branch Grad-CAM++
Model 3 (Proposed Framework)	Territory-Dropout SE-ResNet1D	0.9329	0.7836	0.7849	Multi-Scale Anatomical XAI

5. Qualitative Multi-Scale Attribution & Clinical Visualizations

To evaluate bedside utility, we generated high-resolution multi-panel attribution profiles rendered over centered **3.0-second diagnostic zoom windows** with authentic clinical pink ECG grid lines (0.20 s major grid, 0.04 s minor grid; 0.5 mV major, 0.1 mV minor).

5.1. Acute Myocardial Infarction (MI)

As shown in Fig. 5 (Record #15647, representative ground-truth anterior infarction), Model 3 predicts *MI* with 98.4% confidence. The Meso-level 1D Grad-CAM++ saliency profile highlights the prominent ST-segment elevation and tombstone T-wave inversion across anterior chest leads V2, V3, and V4, while peripheral limb leads remain inactive. At the Macro level, masking the Antero-Septal territory produces a decisive occlusion drop of $\Delta P = 0.340$, with 55.0% of the learned attention weight mapped to the LAD artery. Micro-level Integrated Gradients precisely isolates the J-point elevation.

5.2. ST-T Segment Abnormalities (STTC)

As illustrated in Fig. 6 (Record #10054, representative lateral repolarization abnormality), the Meso-level saliency concentrates over the terminal repolarization phase in lateral leads I, aVL, and V5–V6. The learned attention allocates 47.9% to the Lateral territory (Left Circumflex artery bed), validating the presence of lateral myocardial ischemia.

5.3. Conduction Disturbances (CD)

As depicted in Fig. 7 (Record #4893, representative bundle branch block), for conduction delays, 1D Grad-CAM++ specifically highlights the widened ventricular depolarization interval (QRS duration > 120 ms) across septal leads V1–V3. Territory occlusion sensitivity registers an intense $\Delta P = 0.395$ drop upon masking the Antero-Septal territory, proving that the model bases its decision on intraventricular conduction delays rather than baseline noise.

5.4. Ventricular Hypertrophy (HYP)

In Fig. 8 (Record #16182, representative left ventricular hypertrophy), the model detects Left Ventricular Hypertrophy. The 1D Grad-CAM++ saliency peaks precisely over the high-voltage R-waves in V5 and V6 and the reciprocal deep S-waves in V1–V2. The Lateral territory captures 59.3% of the attention share in this case, demonstrating adherence to the Sokolow-Lyon voltage index ($S_{V1} + R_{V5} > 3.5$ mV).

5.5. Normal Sinus Rhythm Control (NORM)

As shown in Fig. 9 (Record #2083), passing a healthy ECG through the multi-scale engine produces a diffuse, balanced attribution profile across all four vascular territories. Occlusion of any individual lead group yields near-zero probability drop ($\Delta P = 0.101$), confirming that normal predictions reflect genuine global physiological harmony rather than localized artifact reliance.

6. Discussion and Biophysical Interpretability

The empirical findings from this study demonstrate that integrating coronary vascular domain structure directly into 1D deep neural networks resolves the long-standing tension between diagnostic accuracy and clinical interpretability. By partitioning 12-lead ECGs into four anatomical branches and enforcing stochastic territory-dropout during training, Model 3 prevents spatial shortcut co-adaptation, achieving state-of-the-art discrimination on PTB-XL (0.9329 ROC-AUC, 0.7836 F1).

Crucially, our pivot from generative counterfactuals (DDPMs/GANs) to the Multi-Scale Anatomical Explainability Framework eliminates the clinical risk of generative hallucinations. Generative models applied to multi-lead time-series frequently suffer from temporal mode collapse, R-peak destruction, and heart rate drift, synthesizing waveforms that confuse clinicians. In contrast, our multi-scale attribution engine operates strictly on the patient’s factual recording, providing verifiable evidence across three clinically meaningful tiers:

1. **Vascular Bed Localization:** Macro-level territory occlusion sensitivity directly maps diagnostic alerts to the culprit coronary artery (LAD vs. RCA vs. LCx), bridging AI outputs with cardiac catheterization workflows.




Figure 5: Multi-Scale Anatomical Attribution for a Representative Acute Myocardial Infarction Case (MI, Record #15647). (Top-Left) Multi-lead 3.0s zoomed ECG with 1D Grad-CAM++ morphological saliency highlighting ST-segment elevation in leads V2–V4; (Top-Right) Territory Occlusion Sensitivity showing a severe 0.340 drop upon masking the Antero-Septal territory; (Bottom-Left) Learned Cross-Territory Softmax Attention allocating 55.0% to the LAD vascular bed; (Bottom-Right) Axiomatic Integrated Gradients confirming sample completeness.




Figure 6: Multi-Scale Anatomical Attribution for a Representative ST/T-Change Case (STTC, Record #10054). 1D Grad-CAM++ isolates ST-depression and T-wave inversion in Lateral leads I, aVL, V5, V6 (47.9% case attribution share).




Figure 7: Multi-Scale Anatomical Attribution for a Representative Conduction Disturbance Case (CD, Record #4893). 1D Grad-CAM++ isolates the widened QRS complex and slurred intrinsicoid deflection in Antero-Septal leads ($\Delta P = 0.395$).




Figure 8: Multi-Scale Anatomical Attribution for a Representative Ventricular Hypertrophy Case (*HYP*, Record #16182). Grad-CAM++ highlights tall R-wave voltage in Lateral leads V5–V6 and deep S-wave in V1–V2 (59.3% Lateral case share).




Figure 9: Multi-Scale Anatomical Attribution for Normal Sinus Rhythm (*NORM*, Record #2083). Saliency and attention are evenly distributed across all 12 leads and territories, confirming physiological equilibrium.

2. **Morphological Interval Validation:** Meso-level 1D Grad-CAM++ highlights exact waveform intervals (e.g., ST-elevation, QRS widening), allowing clinicians to verify adherence to standardized AHA/ACC electrocardiological criteria.
3. **Sample-Level Completeness:** Micro-level Integrated Gradients guarantees axiomatic mathematical completeness ($|\sum IG - \Delta Score| < 0.02$), ensuring no hidden gradient leakage occurs.

Furthermore, zero-shot external evaluation on 6,877 CPSC2018 records confirmed that territory-dropout regularization prevents models from memorizing single-vendor electrode noise patterns, yielding robust out-of-distribution transferability (0.9049 AUC for Normal Sinus Rhythm, 0.8462 AUC for Conduction Disturbances).

7. Limitations and Future Work

While the proposed framework establishes a new standard for interpretable deep electrocardiology, several limitations should be noted. First, although our quantitative audit across the PTB-XL Fold 10 test cohort established strong electrophysiological alignment, validating these multi-scale visual profiles in a prospective clinical reader study with practicing cardiologists remains an essential next step prior to deployment in hospital intensive care units. Second, while the model excels on isolated primary pathologies, highly complex multi-morbid patients (e.g., concurrent acute anterior infarction, pre-existing left bundle branch block, and severe left ventricular hypertrophy) display competing morphological features across overlapping territories. Developing multi-label hierarchical attribution schemes that decouple simultaneous pathology attributions represents an exciting frontier. Finally, optimizing the Riemann summation of Integrated Gradients using accelerated path approximations will enable real-time, low-latency execution on embedded bedside patient monitors.

8. Conclusion

We presented an Anatomical Territory-Dropout SE-ResNet1D architecture coupled with a hierarchical Multi-Scale Anatomical Explainability (XAI) Framework for 12-lead ECG interpretation. By decomposing the 12 leads into four coronary vascular territories and applying stochastic lead-territory dropout, our classifier achieves state-of-the-art diagnostic discrimination on PTB-XL (Macro ROC-AUC = **0.9329**, Macro F1 = **0.7836**) and demonstrates robust zero-shot external generalization on 6,877 CPSC2018 records (Macro ROC-AUC = **0.7849**). Our multi-scale explainability suite provides transparent, non-hallucinatory decision support by linking predictions to culprit coronary vascular beds, morphological waveform

intervals, and axiomatic voltage attributions on high-resolution 3.0-second clinical pink-grid displays. This work demonstrates that embedding physiological domain structure into deep learning architectures provides a principled foundation for trustworthy artificial intelligence in clinical cardiology.

CRedit Authorship Contribution Statement

Awais Muhammad Lodhi: Conceptualization, Methodology, Software, Formal Analysis, Investigation, Data Curation, Writing – Original Draft, Writing – Review and Editing, Visualization. **Dr. Ghulam Mustafa:** Supervision, Conceptualization, Writing – Review and Editing.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data Availability Statement

The PTB-XL dataset is publicly available at PhysioNet (<https://physionet.org/content/ptb-xl/>). The CPSC2018 dataset is publicly available at the official CPSC 2018 challenge repository.

References

- [1] N. Strodthoff, P. Wagner, T. Schaeffter, W. Samek, Deep learning for ecg analysis: Benchmarks and insights from ptb-xl, *IEEE Journal of Biomedical and Health Informatics* 25 (5) (2020) 1519–1528. doi:10.1109/JBHI.2020.3022989.
- [2] P. Wagner, N. Strodthoff, R.-D. Bousseljot, D. Kreiseler, F. I. Lunze, W. Samek, T. Schaeffter, Ptb-xl, a large publicly available electrocardiography dataset, *Scientific Data* 7 (1) (2020) 154. doi:10.1038/s41597-020-0495-6.
- [3] A. H. Ribeiro, M. H. Ribeiro, G. M. Paixão, D. M. Oliveira, P. R. Gomes, J. A. Canazart, M. P. Ferreira, C. R. Andersson, P. W. Macfarlane, W. Meira Jr, et al., Automatic diagnosis of the 12-lead ecg using a deep neural network, *Nature Communications* 11 (1) (2020) 1760. doi:10.1038/s41467-020-15432-4.
- [4] A. Y. Hannun, P. Rajpurkar, M. Haghpanahi, G. H. Tison, C. Bourn, M. P. Turakhia, A. Y. Ng, Cardiologist-level arrhythmia detection and classification in ambulatory electrocardiograms using a deep neural network, *Nature Medicine* 25 (1) (2019) 65–69. doi:10.1038/s41591-018-0268-3.

- [5] M. Sundararajan, A. Taly, Q. Yan, [Axiomatic attribution for deep networks](#), in: International Conference on Machine Learning, 2017, pp. 3319–3328. URL <http://proceedings.mlr.press/v70/sundararajan17a.html>
- [6] S. M. Lundberg, S.-I. Lee, A unified approach to interpreting model predictions, in: Advances in Neural Information Processing Systems, Vol. 30, 2017. doi:10.48550/arXiv.1705.07874.
- [7] R. R. Selvaraju, M. Cogswell, A. Das, A. Vedaldi, D. Parikh, D. Batra, Grad-cam: Visual explanations from deep networks via gradient-based localization, International Journal of Computer Vision 128 (2) (2020) 336–359. doi:10.1007/s11263-019-01228-7.
- [8] A. A. Ismail, M. Gunady, H. Corrada Bravo, S. Feizi, Benchmarking deep learning interpretability in time series predictions, in: Advances in Neural Information Processing Systems, Vol. 33, 2020, pp. 6441–6452. doi:10.48550/arXiv.2010.13924.
- [9] S. Tonekaboni, S. Joshi, M. D. McCradden, A. Goldenberg, What went wrong and when? instance-wise feature importance for time-series black-box models, in: Advances in Neural Information Processing Systems, Vol. 33, 2020, pp. 7982–7992. doi:10.48550/arXiv.2003.02821.
- [10] O. Korkmaz, Deep learning-based robust ecg generation and classification, Computers in Biology and Medicine 135 (2021) 104646. doi:10.1016/j.combiomed.2021.104646.
- [11] N. Neifar, et al., Diffecg: A versatile probabilistic diffusion model for ecg signals synthesis, in: 2024 IEEE/ACIS 22nd International Conference on Software Engineering Research, Management and Applications (SERA), IEEE, 2024, pp. 182–188. doi:10.48550/arXiv.2306.01875.
- [12] A. Chattopadhyay, A. Sarkar, P. Howlader, V. N. Balasubramanian, Grad-cam++: Generalized gradient-based visual explanations for deep convolutional networks, in: 2018 IEEE Winter Conference on Applications of Computer Vision (WACV), IEEE, 2018, pp. 839–847. doi:10.1109/WACV.2018.00097.
- [13] F. Liu, C. Liu, L. Zhao, X. Zhang, X. Wu, X. Xu, Y. Ding, S. Yin, Z. Chen, J. Li, et al., An open access database for evaluating the algorithms of electrocardiogram rhythm and morphology abnormality detection, Journal of Medical Imaging and Health Informatics 8 (7) (2018) 1368–1373. doi:10.1166/jmhi.2018.2442.
- [14] H. I. Fawaz, G. Forestier, J. Weber, L. Idoumghar, P.-A. Muller, Deep learning for time series classification: a review, Data Mining and Knowledge Discovery 33 (4) (2019) 917–963. doi:10.1007/s10618-019-00619-1.
- [15] K. C. Siontis, P. A. Noseworthy, Z. I. Attia, P. A. Friedman, Artificial intelligence-enhanced electrocardiography in cardiovascular disease management, Nature Reviews Cardiology 18 (7) (2021) 465–478. doi:10.1038/s41569-020-00503-2.
- [16] S. Smigiel, K. Palka, et al., Ecg signal classification using deep learning techniques based on the ptb-xl dataset, Entropy 23 (9) (2021) 1121. doi:10.3390/e23091121.
- [17] T. Mehari, N. Strodthoff, Self-supervised representation learning from 12-lead ecg data, Computers in Biology and Medicine 141 (2022) 105114. doi:10.1016/j.combiomed.2021.105114.
- [18] C. Che, P. Zhang, M. Zhu, Y. Qu, B. Jin, Constrained transformer network for ecg signal processing and arrhythmia classification, BMC Medical Informatics and Decision Making 21 (1) (2021) 1–15. doi:10.1186/s12911-021-01546-2.
- [19] J. Adebayo, J. Gilmer, M. Muelly, I. Goodfellow, M. Hardt, B. Kim, Sanity checks for saliency maps, in: Advances in Neural Information Processing Systems (NeurIPS), Vol. 31, 2018, pp. 9505–9515.